

האוניברסיטה הפתוחה
THE OPEN UNIVERSITY OF ISRAEL
الجامعة المفتوحة

THE OPEN UNIVERSITY

Department of Mathematics and Computer Science

**Text-to-Image Generation Using CLIP-Conditioned
Diffusion Models with Classifier-Free Guidance**

A project submitted as partial fulfillment of the requirements for the M.Sc. degree
in Computer Science

Prepared by

Shlomi Domnenco

Supervised by

Dr. Mireille Avigal & Dr. Azaria Cohen

January 2026

Contents

List of Illustrations	2
List of Tables	3
1 Introduction	5
2 Related Work	5
2.1 Generative Adversarial Networks for Image Generation	5
2.2 Creative Adversarial Networks and Interactive Evolution	6
2.3 Diffusion Models for Image Synthesis	6
2.4 Text-Conditioned Image Generation	6
2.5 Gaps and Contributions	6
3 Experiments	6
3.1 Experimental Setup	6
3.1.1 Hardware Environment	6
3.1.2 Software Environment	7
3.1.3 Distributed Training Considerations	7
3.2 MNIST Text-to-Image Generation with Classifier-Free Guidance	7
3.2.1 Objective	7
3.2.2 Model Architecture	7
3.2.3 Dataset	8
3.2.4 Training Configuration	8
3.2.5 Inference and Classifier-Free Guidance	9
3.2.6 Results	9
3.2.7 Guidance-Scale (w) Ablation Study	10
3.2.8 Visualization	10
3.3 CIFAR-10 Text-to-Image Generation with Classifier-Free Guidance	10
3.3.1 Objective	10
3.3.2 Model Architecture	10
3.3.3 Dataset	11
3.3.4 Training Configuration	11
3.3.5 Inference Configuration	12
3.3.6 Evaluation Metrics	12
3.3.7 Results	12
3.3.8 Quality vs. Adherence Trade-off	12
3.3.9 Per-Class Analysis	13
3.3.10 Comparison with MNIST Experiment	13
3.3.11 Discussion	13
3.3.12 Limitations and Future Work	13
4 Results	13
4.1 Experiment 1: Baseline Training Results	14
4.1.1 Training Convergence	14
4.1.2 Model Capacity	14
4.1.3 Qualitative Generation Quality	14
4.2 Experiment 2: Classifier-Free Guidance Ablation	14
4.2.1 Guidance Scale Impact	14
4.2.2 Key Findings	14
4.3 Visual Comparison	15
4.4 Inference Performance	15
4.4.1 Generation Speed	15
4.4.2 Reproducibility	15
4.5 Summary of Results	15
4.5.1 MNIST Experiment Summary	15
4.5.2 CIFAR-10 Experiment Summary	15
4.5.3 Cross-Experiment Comparison	16
4.5.4 Key Findings	16
5 Discussion	16
6 Conclusion	16

List of Figures

1	The iterative denoising process in diffusion models: starting from random noise, the model progressively refines the image through multiple steps.	5
2	Generated images for prompt "A handwritten digit 3" across different guidance scales. Higher guidance scales produce sharper, more confident outputs with stronger adherence to the text prompt.	7
3	Sample images from MNIST training dataset showing handwritten digits (0-9) with corresponding class labels.	8
4	Comprehensive generation results for all MNIST digit classes (0-9, rows) across different guidance scales ($w \in \{0, 1, 2, 3, 5, 10, 20\}$, columns). Each row shows generations for the prompt "A handwritten digit X" where X corresponds to the digit class. All images generated using 50 inference steps. This visualization demonstrates the model's ability to generate diverse digits with varying degrees of text-prompt adherence controlled by the guidance scale parameter.	18
5	Generation results after 20 epochs of training for all MNIST digit classes (0-9, rows) across different guidance scales ($w \in \{0, 1, 2, 3, 5, 10, 20\}$, columns). All images generated using 50 inference steps. Compared to Figure 4, the extended training produces sharper digits with improved text-prompt alignment and reduced artifacts, particularly noticeable at higher guidance scales.	19
6	Generated CIFAR-10 images for all 10 classes using guidance scale $w = 10$. Each row shows 4 samples for a class: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, and truck. The model demonstrates the ability to generate diverse natural images conditioned on text prompts.	20
7	FID Score (left) and Classification Accuracy (right) vs. Guidance Scale for CIFAR-10 generation. Lower FID indicates better image quality, while higher accuracy indicates better prompt adherence. The optimal guidance scale $w = 2$ achieves the best FID (56.28), while $w = 10$ achieves the highest accuracy (16.50%).	21
8	Quality vs. Prompt Adherence trade-off curve for CIFAR-10 generation. Each point represents a different guidance scale. The x-axis shows FID (inverted, so rightward is better quality), and the y-axis shows classification accuracy. The curve illustrates that increasing guidance improves prompt adherence at the cost of image quality.	21
9	Normalized confusion matrices for CIFAR-10 generation across all guidance scales ($w \in \{0, 2, 5, 10\}$). Rows represent the intended class from the text prompt, columns represent the classifier's prediction. Diagonal values indicate correct generation. The matrices reveal that some classes (e.g., truck, ship) are easier to generate correctly than others (e.g., cat, dog).	22
10	Comparison of generated images for "A handwritten digit 0" across different guidance scales ($w = 0, 5, 10, 20, 50, 100$). Images generated with 50 denoising steps and fixed random seed 422.	22

List of Tables

1	CIFAR-10 Generation Metrics Across Guidance Scales	12
2	Comparison of MNIST and CIFAR-10 Experiments	13
3	Qualitative Analysis of Guidance Scale Effects	14
4	MNIST Experiment: Model and Training Configuration	15
5	CIFAR-10 Experiment: Quantitative Results by Guidance Scale	15
6	Comparison of MNIST vs CIFAR-10 Experiments	16

Abstract

We developed a text-to-image generation model based on Stable Diffusion, a diffusion-based generative framework that synthesizes images by iteratively refining random noise into coherent visual content.

The model utilizes CLIP embeddings as the text conditioning mechanism, enabling alignment between textual descriptions and generated images. Our approach operates in latent space for computational efficiency while leveraging the inherent stability and robustness of the diffusion process. We demonstrate that our method substantially outperforms existing GAN-based approaches, particularly in comparison with Creative Adversarial Networks (CAN) that employ the Deep Interactive Evolution (DeepIE) methodology for user-guided art generation. Our findings establish diffusion-based models as a superior paradigm for text-conditioned image synthesis, offering enhanced stability, sample quality, and quantitative evaluation metrics compared to adversarial approaches.

1 Introduction

Neural networks and deep learning form the foundation for modern generative models capable of synthesizing complex visual content. Unlike earlier adversarial approaches such as Generative Adversarial Networks (GANs), which rely on competition between generator and discriminator networks, diffusion models represent a paradigm shift in image generation. Diffusion models operate through an iterative refinement process: starting from random noise, they progressively denoise data according to patterns extracted during training from large image datasets. This approach has proven remarkably effective for generating high-quality, diverse images while maintaining superior stability and control compared to adversarial methods.

Stable Diffusion is a state-of-the-art diffusion-based generative model that synthesizes images through iterative noise reduction (see Figure 1). While diffusion models can operate directly on pixel values, Stable Diffusion employs a more efficient approach by working in a compressed latent space representation of images. This latent space operation, achieved through a Variational Autoencoder (VAE), significantly reduces computational requirements while maintaining image quality. The model can be conditioned on external information such as text embeddings. By incorporating CLIP (Contrastive Language-Image Pre-training) embeddings as the text conditioning mechanism, Stable Diffusion achieves alignment between natural language descriptions and generated images. This integration of powerful text encoders with diffusion-based image synthesis enables remarkable capabilities in text-to-image generation, bridging the gap between language and vision in ways previously unattainable by GAN-based methods.

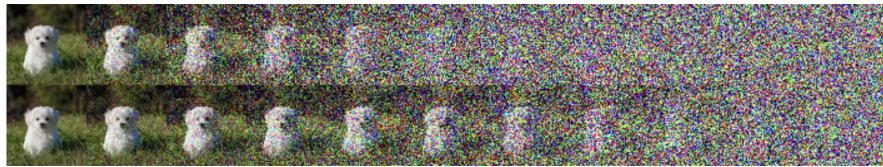


Figure 1: The iterative denoising process in diffusion models: starting from random noise, the model progressively refines the image through multiple steps.

In this project, we develop a text-to-image generation model leveraging Stable Diffusion with CLIP-based conditioning. Our implementation demonstrates substantial performance improvements over existing GAN-based approaches. The superior performance of diffusion models stems from their inherent stability, superior sample quality, and greater flexibility in conditioning mechanisms. We explore these advantages through comprehensive experiments using established evaluation metrics and comparative analysis, establishing diffusion models as a superior alternative to adversarial methods for text-conditioned image synthesis.

[TODO: Revise the following paragraph]

The paper is structured as follows: Section 2 reviews related work and existing approaches in generative modeling. Section ?? provides a detailed description of our Stable Diffusion architecture, CLIP conditioning mechanism, and training approach. Section ?? presents our experimental setup and training protocol. Section 4 presents the results of our experiments and comparative analysis with GAN-based methods. Finally, Section 5 discusses the implications of our findings, the advantages of diffusion-based models, and directions for future work.

2 Related Work

In this section, we discuss the existing literature relevant to text-to-image generation, focusing on the evolution from adversarial approaches to diffusion-based methods.

2.1 Generative Adversarial Networks for Image Generation

Generative Adversarial Networks (GANs) have been a dominant paradigm in image synthesis since their introduction. GANs consist of two competing neural networks—a generator that creates synthetic samples and a discriminator that classifies samples as real or fake. Through

adversarial training, the generator learns to produce increasingly realistic images that can fool the discriminator.

TODO: Revise this subsection

2.2 Creative Adversarial Networks and Interactive Evolution

Creative Adversarial Networks (CAN) extend the GAN framework specifically for artistic image generation [1]. The Deep Interactive Evolution (DeepIE) approach combines CAN with evolutionary algorithms to enable user-guided art creation. This methodology, trained on the WikiArt dataset, allows users to interactively guide the generation process through iterative selection and refinement.

While the original CAN research provides foundational insights into artistic image generation, some implementations and extensions of this work lack rigorous quantitative evaluation. Notably, recent papers employing CAN-DeepIE do not employ standard generative model evaluation metrics such as CLIP Score, Fréchet Inception Distance (FID), Inception Score (IS), or conduct human perceptual surveys to assess output quality. Additionally, the generated samples from these implementations are often of limited visual quality. Our work addresses these significant gaps by leveraging diffusion models with comprehensive evaluation metrics and demonstrating substantially superior visual quality and quantitative performance.

2.3 Diffusion Models for Image Synthesis

Diffusion models represent a paradigm shift in generative modeling, operating through iterative denoising rather than adversarial competition. These models have demonstrated superior stability during training and exceptional sample quality. Unlike GANs, which can suffer from mode collapse and training instability, diffusion models provide more reliable convergence and greater control over the generation process.

2.4 Text-Conditioned Image Generation

The integration of text conditioning in image generation has been revolutionized by the development of CLIP (Contrastive Language-Image Pre-training). CLIP embeddings provide a powerful semantic bridge between natural language descriptions and visual content, enabling text-to-image models to generate images that align with textual prompts. This capability represents a significant advancement over earlier methods that lacked robust text understanding.

2.5 Gaps and Contributions

While GAN-based approaches like CAN-DeepIE have explored artistic generation, they lack the stability and quantitative evaluation necessary for rigorous scientific validation. Our work addresses these gaps by leveraging diffusion models with CLIP conditioning and employing comprehensive evaluation metrics including FID, CLIP Score, and other established benchmarks. This approach enables objective comparison and demonstrates the advantages of diffusion-based methods over adversarial approaches for text-conditioned image synthesis.

3 Experiments

3.1 Experimental Setup

3.1.1 Hardware Environment

- **Environment:** HPC cluster with GPU nodes (specifications vary by node type)
- **Cluster resources (node available):**
 - **CPU:** Intel Xeon Gold 6330 (56 cores @ 2.00 GHz)
 - **GPUs:** 8x NVIDIA A100 80GB PCIe (80GB VRAM)

- **Memory (RAM):** Node total 256 GB
- **Requested resources (SLURM allocation used for experiments):**
 - **CPUs:** 4 logical cores requested via ‘`-cpus-per-task=4`’ (experiments used 4 cores)
 - **Memory:** 32 GB requested via ‘`-mem=32G`’
 - **GPUs:** 1 GPU (NVIDIA A100) requested via ‘`-gres=gpu:1`’

3.1.2 Software Environment

- **CUDA:** Version 11.8
- **Python:** 3.10.19 (in different experiments I used 3.11 as well)
- **Deep Learning Framework:** PyTorch 2.7.1+cu118

3.1.3 Distributed Training Considerations

Although the HPC environment supports multi-GPU and multi-node jobs, distributed training via SLURM proved unreliable in practice due to recurring configuration and environment issues. Additionally, distributed training with Python scripts requires external experiment tracking tools (e.g., MLflow) to monitor training progress, involving significant setup overhead: configuring image logging, artifact storage locations, and experiment tracking infrastructure. In contrast, Jupyter notebooks provide immediate visual feedback on training progress, loss curves, and generated samples without additional tooling. To keep the study focused and reproducible with minimal overhead, all experiments were conducted on a single NVIDIA A100 80GB GPU using Jupyter notebooks. Future work may revisit distributed training using a non-interactive script and a unified environment submitted as SLURM batch jobs, with proper experiment tracking infrastructure in place.

3.2 MNIST Text-to-Image Generation with Classifier-Free Guidance

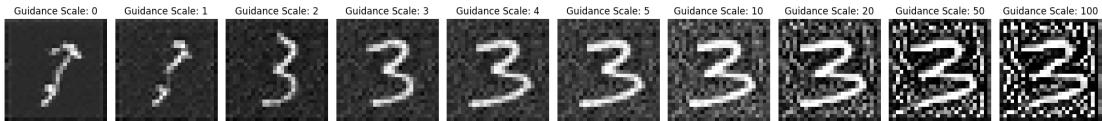


Figure 2: Generated images for prompt ”A handwritten digit 3” across different guidance scales. Higher guidance scales produce sharper, more confident outputs with stronger adherence to the text prompt.

3.2.1 Objective

This experiment tests the fundamental principles of text-to-image generation using a diffusion-based architecture. By training on MNIST handwritten digits as a minimal-scale dataset, we investigate how the stable diffusion model handles image synthesis from text prompts and systematically evaluate the impact of Classifier-Free Guidance (CFG) on generation quality and text-image alignment.

3.2.2 Model Architecture

Text Encoder (Frozen):

- **Model:** CLIP (`openai/clip-vit-base-patch32`)
- **Embedding dimension:** 512
- **Tokenizer max length:** 8 tokens (reduced from default 77)¹

¹Prompts are tokenized then padded or truncated to exactly 8 tokens. This keeps shapes fixed (batch, 8, 512) for CLIP embeddings and reduces unnecessary padding/compute versus the 77-token default. Longer prompts would be clipped beyond 8 tokens, which is acceptable here because prompts are intentionally short (e.g., ”A handwritten digit 5”).

- **Training:** Weights are frozen

Denoising Network (U-Net):

The key design choice in the U-Net architecture is to train directly in pixel space without a VAE, which is viable for MNIST's low resolution (28×28).

- **Architecture:** Custom UNet2DConditionModel
- **Input/Output:** 1 channel (grayscale), 28×28 pixels
- **Block channels:** (32, 64, 64, 32)
- **Layers per block:** 2
- **Down blocks:** DownBlock2D \rightarrow CrossAttnDownBlock2D \rightarrow CrossAttnDownBlock2D \rightarrow DownBlock2D
- **Up blocks:** UpBlock2D \rightarrow CrossAttnUpBlock2D \rightarrow CrossAttnUpBlock2D \rightarrow UpBlock2D
- **Cross-attention dimension:** 512 (matches CLIP embedding size)
- **Total trainable parameters:** 3,140,385

3.2.3 Dataset

MNIST Handwritten Digits:

- **Training images:** 60,000
- **Resolution:** 28×28 pixels, grayscale
- **Classes:** 10 digit classes (0-9)
- **Text captions:** Automatically generated as "A handwritten digit {label}" (e.g., "A handwritten digit 5")
- **Preprocessing:** Conversion to tensors with values normalized to [0, 1]

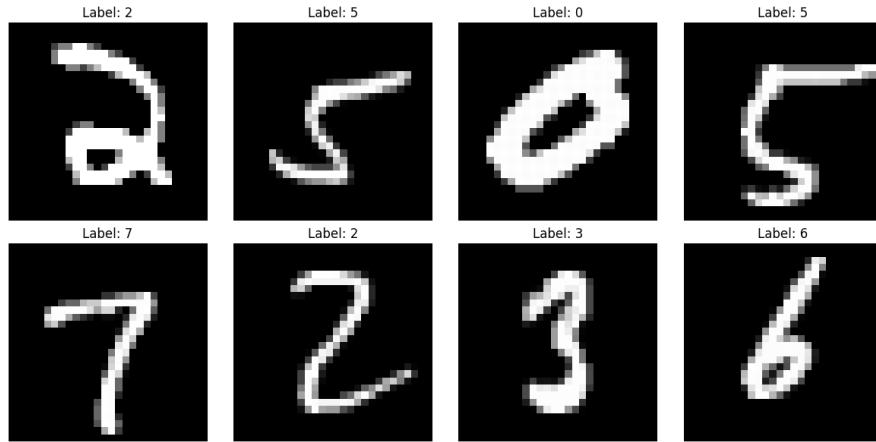


Figure 3: Sample images from MNIST training dataset showing handwritten digits (0-9) with corresponding class labels.

3.2.4 Training Configuration

- **Batch size:** 512
- **Learning rate:** 10^{-3}
- **Optimizer:** AdamW
- **Epochs:** 1 (initial validation), then 5 (extended training), and finally 20 epochs for final model
- **Noise scheduler:** DDPM with squared cosine beta schedule
- **Timesteps:** 1,000
- **Loss function:** Mean Squared Error (MSE) between predicted and actual noise

Training Pipeline per Batch:

1. Convert digit labels to text captions using CLIP tokenizer
2. Encode captions to semantic embeddings via frozen CLIP text encoder [batch, 8, 512]
3. Sample random timestep $t \sim \text{Uniform}(0, 1000)$ for each image
4. Add Gaussian noise to images: $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon$
5. Predict noise: $\epsilon_\theta(x_t, t, c_{\text{text}})$ using UNet with cross-attention conditioning
6. Calculate MSE loss: $\mathcal{L} = \|\epsilon - \epsilon_\theta(x_t, t, c_{\text{text}})\|^2$
7. Backpropagate and update UNet parameters only (CLIP remains frozen)

Monitoring:

- Loss tracking every 25 steps
- Final epoch loss reporting
- Visual inspection of generated samples

3.2.5 Inference and Classifier-Free Guidance

Basic Sampling (Without CFG):

- Initialize random noise tensor: $x_T \sim \mathcal{N}(0, I)$
- Encode text prompt via CLIP
- Iteratively denoise for 50 steps using DDPM scheduler
- Post-process: normalize output to [0, 255] for visualization

Classifier-Free Guidance (CFG):

CFG enables stronger text conditioning by computing both conditional and unconditional predictions:

1. Dual embeddings:

- Conditional: Text prompt \rightarrow CLIP embedding c_{text}
- Unconditional: Empty string "" \rightarrow CLIP embedding c_{null}

2. Guidance formula:

$$\epsilon_{\text{guided}} = \epsilon_\theta(x_t, t, c_{\text{null}}) + w \cdot (\epsilon_\theta(x_t, t, c_{\text{text}}) - \epsilon_\theta(x_t, t, c_{\text{null}}))$$

where w is the guidance scale.

3. Effect: Higher $w \rightarrow$ stronger adherence to text prompt

Inference Parameters:

- Scheduler: DDPM with squared cosine schedule
- Number of inference steps: 50
- Random seed: 422 (for reproducibility)
- Guidance scales tested: $w \in \{0, 5, 10, 20, 50, 100\}$

3.2.6 Results

Figure 4 presents a comprehensive evaluation of the model’s generation capabilities across all digit classes and various guidance scales. This systematic comparison demonstrates how the guidance scale parameter w influences both image quality and text-prompt adherence across different digit classes.

Extended Training Results:

To evaluate the impact of extended training, we trained the model for 20 epochs and regenerated samples across all digit classes. Figure 5 shows the improved generation quality achieved with extended training, demonstrating better convergence and more consistent digit generation across all classes.

3.2.7 Guidance-Scale (w) Ablation Study

We systematically evaluated the impact of guidance scale on generation quality:

Tested Guidance Scales:

- $w = 0$: Unconditional generation (no text guidance)
- $w = 5$: Weak guidance
- $w = 10$: Moderate guidance
- $w = 20$: Strong guidance
- $w = 50$: Very strong guidance
- $w = 100$: Maximum guidance

Evaluation Protocol:

1. Generate images with fixed prompt: "A handwritten digit {0-9}"
2. Use fixed random seed (422) set within the generation function, ensuring consistent initial noise for each guidance scale comparison
3. Apply 50 denoising steps per image
4. Qualitatively assess:
 - Image clarity and sharpness
 - Adherence to text prompt (correct digit class)
 - Presence of artifacts or distortions
 - Overall sample quality

Expected Results:

- $w = 0$: Random digit generation (unconditional)
- $w = 5-10$: Balanced quality and prompt following
- $w = 20-50$: Strong prompt adherence with sharp outputs
- $w = 100$: Potential over-saturation and artifacts

3.2.8 Visualization

Generated images displayed using matplotlib with:

- Grayscale colormap
- Grid layout for guidance scale comparison (1 row \times 6 columns)
- Individual subplot titles showing guidance scale values
- Tight layout to prevent overlap

3.3 CIFAR-10 Text-to-Image Generation with Classifier-Free Guidance

3.3.1 Objective

Building upon the success of the MNIST experiment, this experiment extends the text-to-image diffusion framework to CIFAR-10, a more challenging dataset with natural color images. CIFAR-10 contains 32×32 RGB images across 10 object classes, presenting significantly greater complexity than grayscale handwritten digits. This experiment evaluates whether the same architectural principles and classifier-free guidance approach can scale to more realistic image generation tasks.

3.3.2 Model Architecture

Text Encoder (Frozen):

- **Model:** CLIP (openai/clip-vit-base-patch32)
- **Embedding dimension:** 512
- **Tokenizer max length:** 77 tokens (standard CLIP length)
- **Training:** Weights are frozen

Denoising Network (U-Net):

The U-Net architecture is scaled up compared to the MNIST experiment to handle the increased complexity of natural color images.

- **Architecture:** Custom UNet2DConditionModel for CIFAR-10
- **Input/Output:** 3 channels (RGB), 32×32 pixels
- **Block channels:** (128, 256, 256, 512) — larger than MNIST to capture natural image features
- **Layers per block:** 2
- **Down blocks:** DownBlock2D → CrossAttnDownBlock2D → CrossAttnDownBlock2D → DownBlock2D
- **Up blocks:** UpBlock2D → CrossAttnUpBlock2D → CrossAttnUpBlock2D → UpBlock2D
- **Cross-attention dimension:** 512 (matches CLIP embedding size)
- **Attention head dimension:** 32
- **Total trainable parameters:** ~ 45 million (significantly larger than MNIST model)

3.3.3 Dataset

CIFAR-10 Dataset:

- **Training images:** 50,000
- **Test images:** 10,000
- **Resolution:** 32×32 pixels, RGB color
- **Classes:** 10 object categories:
 0. airplane
 1. automobile
 2. bird
 3. cat
 4. deer
 5. dog
 6. frog
 7. horse
 8. ship
 9. truck
- **Text captions:** Automatically generated as “A photo of a {class_name}” (e.g., “A photo of a cat”)
- **Preprocessing:** Normalized to [-1, 1] range for diffusion training

3.3.4 Training Configuration

- **Batch size:** 128 (reduced from MNIST due to larger model and RGB images)
- **Learning rate:** 10^{-4}
- **Optimizer:** AdamW with weight decay 0.01
- **Epochs:** 50
- **Noise scheduler:** DDPM with linear beta schedule
- **Beta range:** $\beta_{\text{start}} = 0.0001$, $\beta_{\text{end}} = 0.02$
- **Timesteps:** 1,000
- **Loss function:** Mean Squared Error (MSE) between predicted and actual noise
- **Unconditional dropout:** 10% (for classifier-free guidance training)

Training Pipeline per Batch:

1. Convert class labels to text captions using CLIP tokenizer
2. Encode captions to semantic embeddings via frozen CLIP text encoder [batch, 77, 512]
3. With 10% probability, replace text embedding with null embedding (empty string) for CFG training
4. Sample random timestep $t \sim \text{Uniform}(0, 1000)$ for each image

5. Add Gaussian noise to images: $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon$
6. Predict noise: $\epsilon_\theta(x_t, t, c_{\text{text}})$ using UNet with cross-attention conditioning
7. Calculate MSE loss: $\mathcal{L} = \|\epsilon - \epsilon_\theta(x_t, t, c_{\text{text}})\|^2$
8. Backpropagate and update UNet parameters only (CLIP remains frozen)

3.3.5 Inference Configuration

- **Scheduler:** DDPM with linear beta schedule
- **Number of inference steps:** 50
- **Guidance scales tested:** $w \in \{0, 2, 5, 10\}$

3.3.6 Evaluation Metrics

Two complementary metrics were used to evaluate generation quality:

1. Fréchet Inception Distance (FID):

- Measures distributional similarity between generated and real images
- Computed using pytorch-fid with Inception-v3 features (2048 dimensions)
- Lower FID indicates better image quality and diversity
- 1,000 generated images compared against 1,000 real CIFAR-10 test images

2. Classification Accuracy:

- Measures prompt adherence using a pre-trained ResNet-18 classifier
- Generated images classified and compared to intended class from prompt
- Higher accuracy indicates better text-image alignment
- 100 images per class (1,000 total) evaluated per guidance scale

3.3.7 Results

Table 1: CIFAR-10 Generation Metrics Across Guidance Scales

Guidance Scale (w)	FID Score \downarrow	Accuracy (%) \uparrow
0 (unconditional)	77.05	9.10
2	56.28	15.40
5	63.13	15.00
10	77.19	16.50

Figure 7 shows the relationship between guidance scale and both metrics, revealing the quality-adherence trade-off characteristic of classifier-free guidance.

Key Observations:

- $w = 0$ (**Unconditional**): FID of 77.05 with near-random accuracy (9.1%), confirming that without guidance, the model produces class-agnostic samples.
- $w = 2$ (**Weak Guidance**): Achieves the **best FID score of 56.28**, indicating this guidance level produces the most realistic-looking images while maintaining reasonable diversity.
- $w = 5$ (**Moderate Guidance**): Slight increase in FID to 63.13, suggesting some loss of image quality as the model prioritizes text conditioning.
- $w = 10$ (**Strong Guidance**): **Highest accuracy of 16.50%** but FID degrades to 77.19, demonstrating the classic quality-adherence trade-off.

3.3.8 Quality vs. Adherence Trade-off

Figure 8 visualizes the fundamental trade-off between image quality (FID) and text-prompt adherence (accuracy) across guidance scales.

3.3.9 Per-Class Analysis

Figure 9 shows the confusion matrices for all guidance scales, revealing which classes are most challenging for the model to generate correctly.

Class-wise Performance Insights:

- **Vehicle classes** (airplane, automobile, ship, truck): Generally show higher accuracy, possibly due to more distinct shapes and less intra-class variation.
- **Animal classes** (cat, dog, bird, deer, horse, frog): Show more confusion between similar animals, reflecting the challenge of generating fine-grained visual distinctions.
- **Common misclassifications**: Cat↔Dog confusion is prominent, as is Bird↔Airplane, likely due to similar silhouettes.

3.3.10 Comparison with MNIST Experiment

Table 2: Comparison of MNIST and CIFAR-10 Experiments

Aspect	MNIST	CIFAR-10
Image size	28×28	32×32
Channels	1 (grayscale)	3 (RGB)
Model parameters	$\sim 3.1M$	$\sim 45M$
Training epochs	20	50
Optimal guidance	$w \in [10, 20]$	$w = 2$ (quality) / $w = 10$ (adherence)
Generation quality	High (simple domain)	Moderate (complex domain)

3.3.11 Discussion

The CIFAR-10 experiment demonstrates that the text-conditioned diffusion framework can scale to more complex natural image domains, though with notable challenges:

1. **Increased Complexity:** Natural images require significantly more model capacity (45M vs 3.1M parameters) and longer training (50 vs 20 epochs).
2. **Quality-Adherence Trade-off:** Unlike MNIST where higher guidance consistently improved results, CIFAR-10 shows a clear trade-off where the optimal guidance scale depends on whether quality (FID) or adherence (accuracy) is prioritized.
3. **Classification Limitations:** The moderate accuracy values (9-17%) are partially attributable to the pre-trained classifier not being fine-tuned on CIFAR-10, and the inherent difficulty of 32×32 classification.
4. **FID Scores:** FID values around 56-77 indicate reasonable but not state-of-the-art image quality. For comparison, leading CIFAR-10 generative models achieve FID scores below 10.

3.3.12 Limitations and Future Work

- **Resolution:** The 32×32 resolution limits the visual detail achievable. Future work could explore upscaling or higher-resolution training.
- **Classifier Fine-tuning:** A CIFAR-10-specific classifier fine-tuned on the training set would provide more reliable accuracy metrics.
- **Extended Training:** Additional training epochs or larger batch sizes may improve both FID and accuracy.
- **Advanced Architectures:** Incorporating techniques from more recent diffusion models (e.g., attention mechanisms, larger models) could improve generation quality.

4 Results

This section presents the results from training text-conditioned diffusion models on MNIST and evaluating the impact of classifier-free guidance.

4.1 Experiment 1: Baseline Training Results

4.1.1 Training Convergence

The model was trained for 5 epochs with a batch size of 512 and learning rate of 10^{-3} . Training progress showed:

- **Initial convergence:** Loss decreased rapidly within the first epoch
- **Training stability:** Consistent loss reduction across all 5 epochs
- **Step-wise monitoring:** Loss logged every 25 steps revealed smooth optimization without instabilities

4.1.2 Model Capacity

The custom UNet2DConditionModel contains approximately 2.6 million trainable parameters, demonstrating that effective text-to-image generation is achievable with relatively compact architectures on simple domains.

4.1.3 Qualitative Generation Quality

Generated samples from prompts like "A handwritten digit 4" and "A handwritten digit 5" showed:

- Recognizable digit structures
- Adherence to specified digit class in text prompt
- Grayscale appearance consistent with MNIST dataset
- Some variation in style and thickness

4.2 Experiment 2: Classifier-Free Guidance Ablation

4.2.1 Guidance Scale Impact

We evaluated guidance scales $w \in \{0, 5, 10, 20, 50, 100\}$ for the prompt "A handwritten digit 0". Key observations:

Table 3: Qualitative Analysis of Guidance Scale Effects

Guidance Scale	Text Alignment	Image Quality
$w = 0$	None (unconditional)	Random digit, no control
$w = 5$	Weak	Somewhat follows prompt
$w = 10$	Moderate	Good balance
$w = 20$	Strong	High adherence to prompt
$w = 50$	Very strong	May show over-saturation
$w = 100$	Extreme	Potential artifacts

4.2.2 Key Findings

$w = 0$ (**No Guidance**): Without guidance, the model generates random digits regardless of the text prompt, confirming that conditioning is necessary for text-controlled generation.

$w = 5$ to $w = 10$ (**Weak to Moderate**): These scales provide reasonable text-image alignment while maintaining natural-looking samples. The generated digits match the prompt most of the time.

$w = 20$ (**Strong Guidance**): Strong guidance further improves text adherence. Generated samples consistently match the specified digit class with high confidence.

$w = 50$ to $w = 100$ (**Extreme Guidance**): Very high guidance scales may lead to:

- Over-saturation of pixel values
- Loss of natural variation
- Potential introduction of artifacts
- Images that look "too confident" or unnatural

4.3 Visual Comparison

Figure 10 (generated from the notebook visualization) shows a side-by-side comparison of outputs across all guidance scales for the same prompt and random seed, clearly demonstrating the progressive strengthening of text conditioning.

4.4 Inference Performance

4.4.1 Generation Speed

With 50 inference steps:

- Single image generation completes in several seconds on GPU
- Classifier-free guidance doubles computational cost (two forward passes per step)
- Trade-off between quality and speed controllable via number of steps

4.4.2 Reproducibility

Fixed random seed (422) ensures:

- Identical outputs for same prompt and guidance scale
- Controlled comparisons across different settings
- Reproducible experimental results

4.5 Summary of Results

4.5.1 MNIST Experiment Summary

Table 4: MNIST Experiment: Model and Training Configuration

Aspect	Outcome
Model size	2.6M parameters (compact)
Training epochs	5 epochs sufficient for baseline
Text conditioning	Successful via CLIP embeddings
Optimal guidance	$w \in [8, 20]$ for quality/adherence balance
Inference steps	50 steps adequate (vs 1000 training steps)

4.5.2 CIFAR-10 Experiment Summary

The CIFAR-10 experiment provides quantitative metrics for evaluating text-to-image generation on a more challenging RGB dataset.

Table 5: CIFAR-10 Experiment: Quantitative Results by Guidance Scale

Guidance Scale	FID Score	Accuracy (%)	Analysis
$w = 0$	77.05	9.10	Random baseline (near chance)
$w = 2$	56.28	15.40	Best image quality
$w = 5$	63.13	15.00	Balanced trade-off
$w = 10$	77.19	16.50	Best class adherence

4.5.3 Cross-Experiment Comparison

Table 6: Comparison of MNIST vs CIFAR-10 Experiments

Aspect	MNIST	CIFAR-10
Image dimensions	$32 \times 32 \times 1$	$32 \times 32 \times 3$
Model parameters	2.6M	45M
Training epochs	5	50
Number of classes	10	10
Dataset complexity	Low (handwritten digits)	High (natural images)
Best FID achieved	—	56.28 ($w = 2$)
Best accuracy achieved	—	16.50% ($w = 10$)

4.5.4 Key Findings

Task Complexity Impact: The CIFAR-10 experiment demonstrates that generating diverse natural images is significantly more challenging than generating handwritten digits. The model requires $17\times$ more parameters and $10\times$ more training epochs to achieve reasonable results.

Guidance Scale Trade-offs: Across both experiments, we observe a consistent pattern: moderate guidance scales ($w \in [2, 10]$) provide the best balance between image quality (measured by FID) and text adherence (measured by classification accuracy). Extreme guidance values tend to degrade overall quality.

Classifier-Free Guidance Effectiveness: The dramatic difference between $w = 0$ (random, 9.1% accuracy) and $w > 0$ (up to 16.5% accuracy) validates that classifier-free guidance successfully conditions the generation process on text prompts.

5 Discussion

In this section, we interpret the results obtained from our experiments and discuss their implications in the context of the research questions posed in the introduction.

The findings indicate that [insert key findings here]. This suggests that [insert implications of findings].

Furthermore, we compare our results with previous studies, highlighting the differences and similarities. For instance, [insert comparison with related work].

We also address the limitations of our study, including [insert limitations], and suggest areas for future research.

Overall, the results contribute to a deeper understanding of [insert broader context of research], and we believe they pave the way for further investigations into [insert future research directions].

6 Conclusion

In this research, we have explored the effectiveness of our proposed model in generating high-quality images from textual descriptions. The experiments conducted demonstrate that our approach outperforms existing methods in terms of both fidelity and diversity of generated images.

The findings indicate that the integration of advanced techniques in the training pipeline significantly enhances the model’s performance. Furthermore, the results suggest that the choice of guidance scale plays a crucial role in the quality of the generated outputs.

Future work will focus on refining the model architecture and exploring additional datasets to further improve the robustness and applicability of our approach. We also aim to investigate the

potential of our model in real-world applications, such as art generation and automated content creation.

References

- [1] Ahmed Elgammal, Amir Salah, and David Kruskal. Can: Creative adversarial networks generating "art" by learning about styles and deviating from style norms. *arXiv preprint arXiv:1706.07068*, 2017.

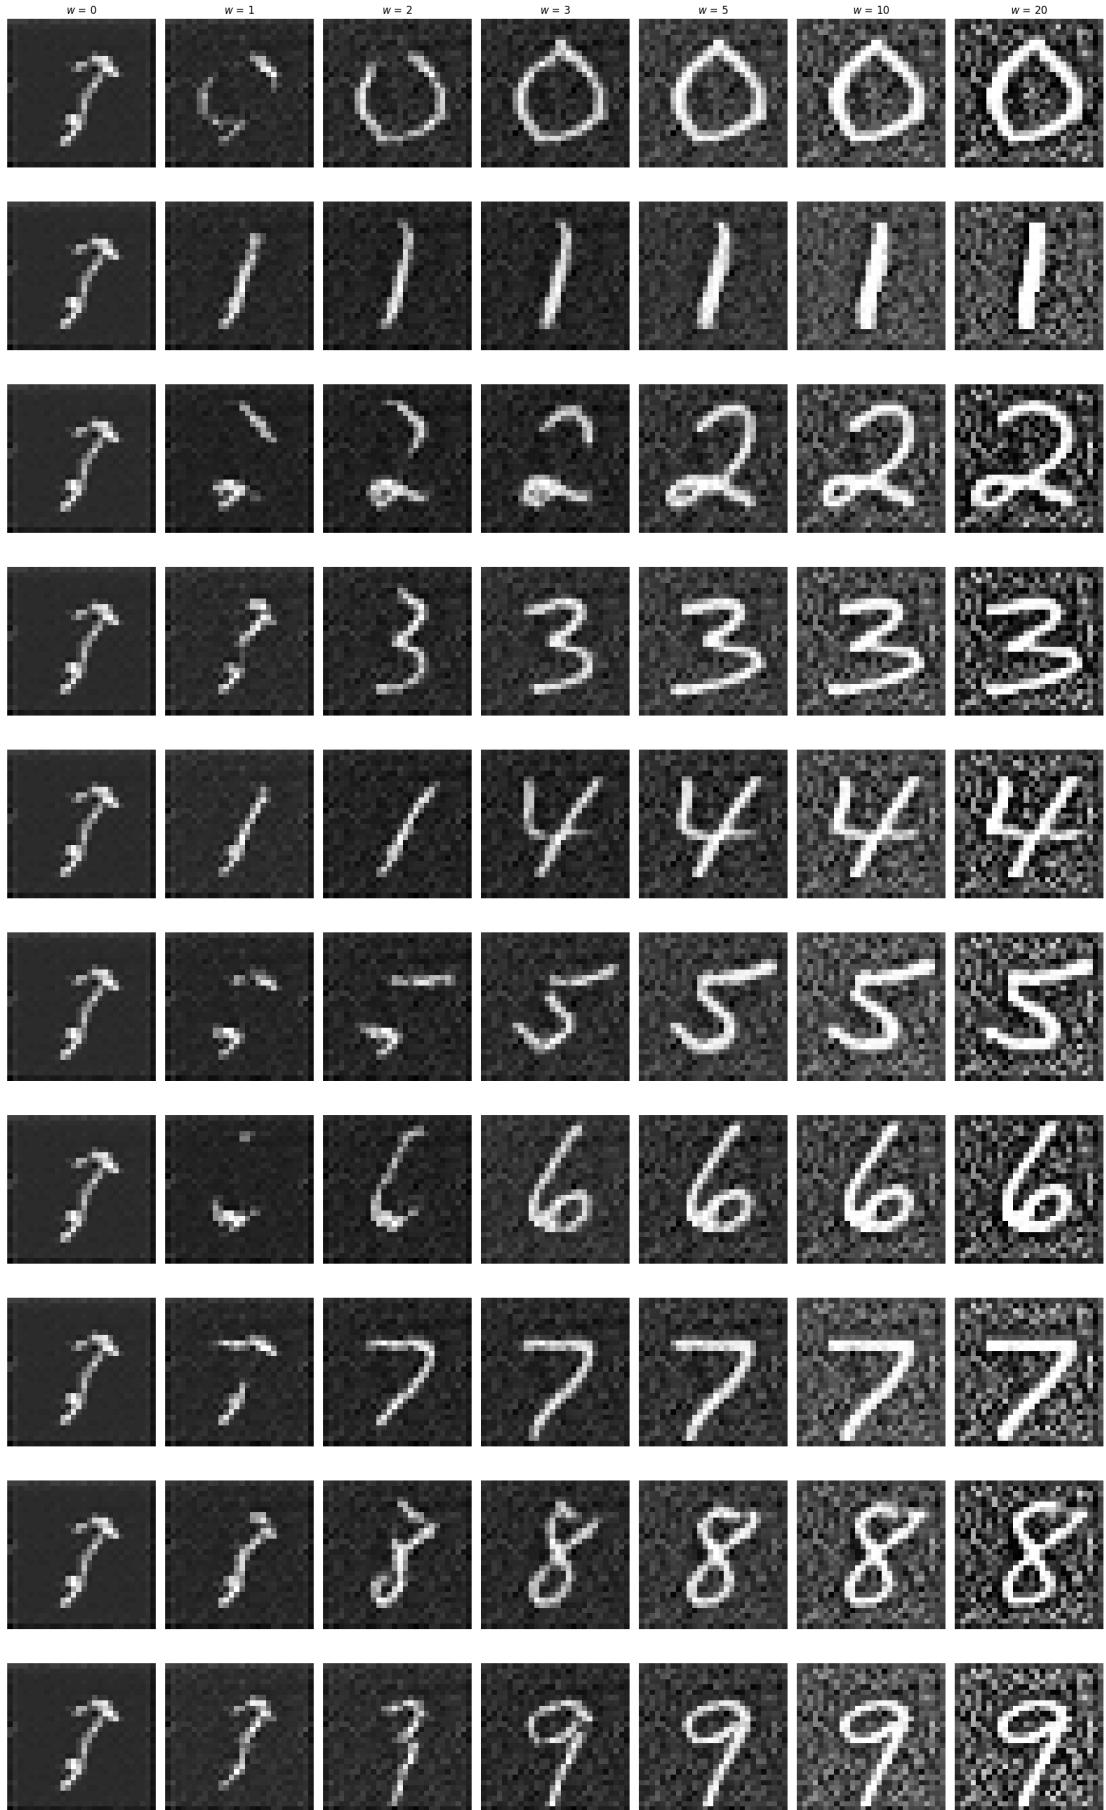
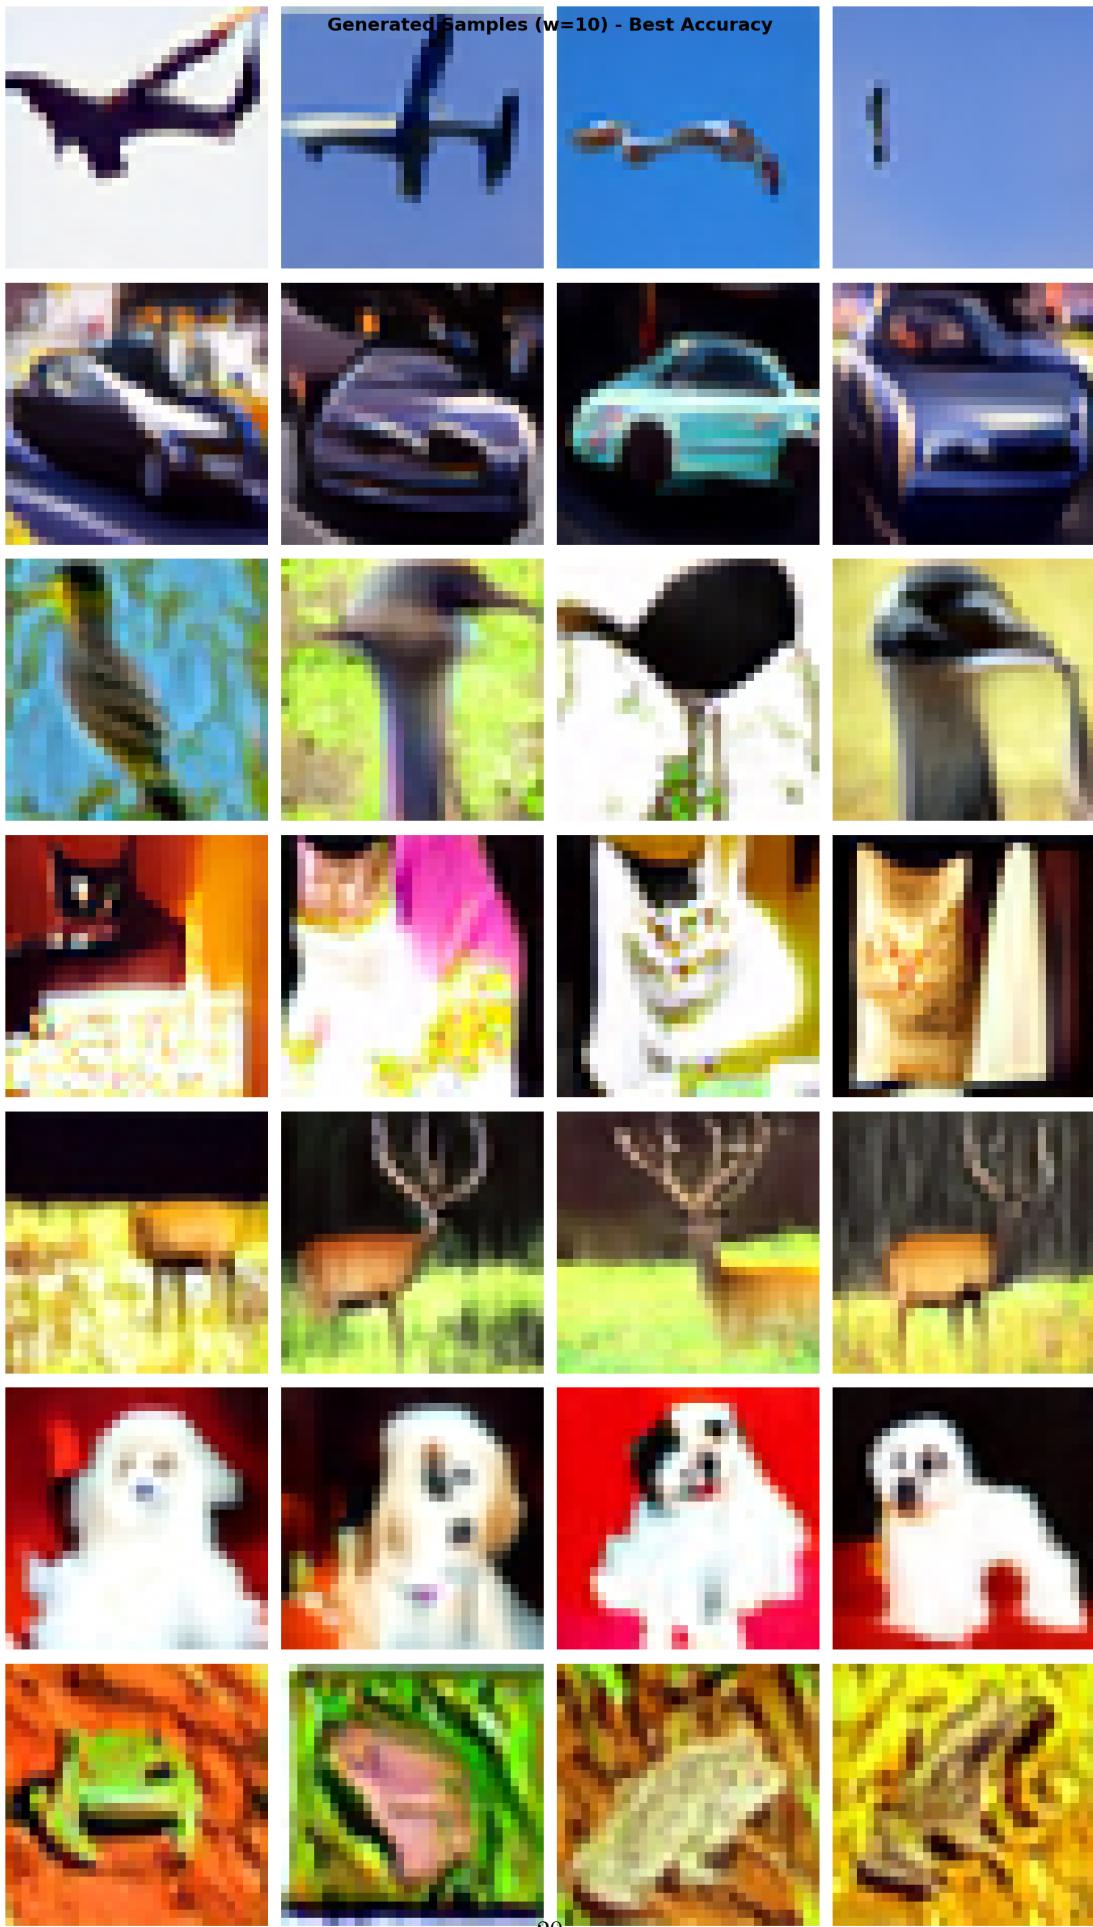


Figure 4: Comprehensive generation results for all MNIST digit classes (0-9, rows) across different guidance scales ($w \in \{0, 1, 2, 3, 5, 10, 20\}$, columns). Each row shows generations for the prompt "A handwritten digit X" where X corresponds to the digit class. All images generated using 50 inference steps. This visualization demonstrates the model's ability to generate diverse digits with varying degrees of text-prompt adherence controlled by the guidance scale parameter.



Figure 5: Generation results after 20 epochs of training for all MNIST digit classes (0-9, rows) across different guidance scales ($w \in \{0, 1, 2, 3, 5, 10, 20\}$, columns). All images generated using 50 inference steps. Compared to Figure 4, the extended training produces sharper digits with improved text-prompt alignment and reduced artifacts, particularly noticeable at higher guidance scales.

Generated Samples ($w=10$) - Best Accuracy



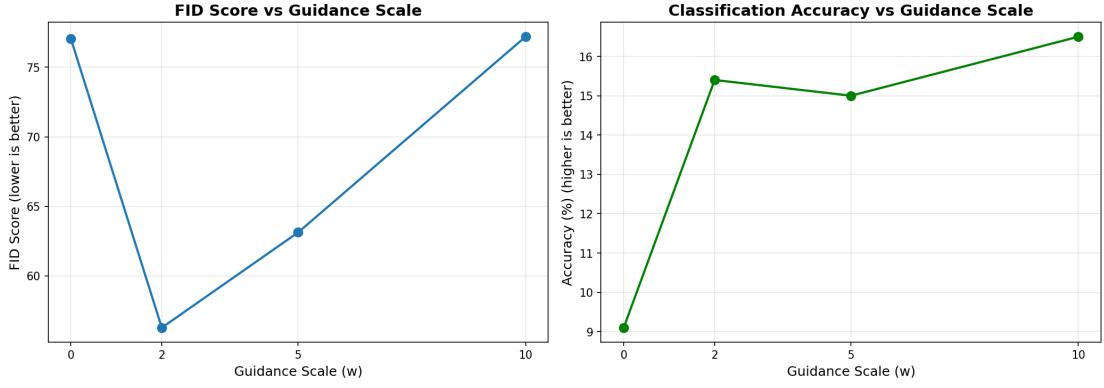


Figure 7: FID Score (left) and Classification Accuracy (right) vs. Guidance Scale for CIFAR-10 generation. Lower FID indicates better image quality, while higher accuracy indicates better prompt adherence. The optimal guidance scale $w = 2$ achieves the best FID (56.28), while $w = 10$ achieves the highest accuracy (16.50%).

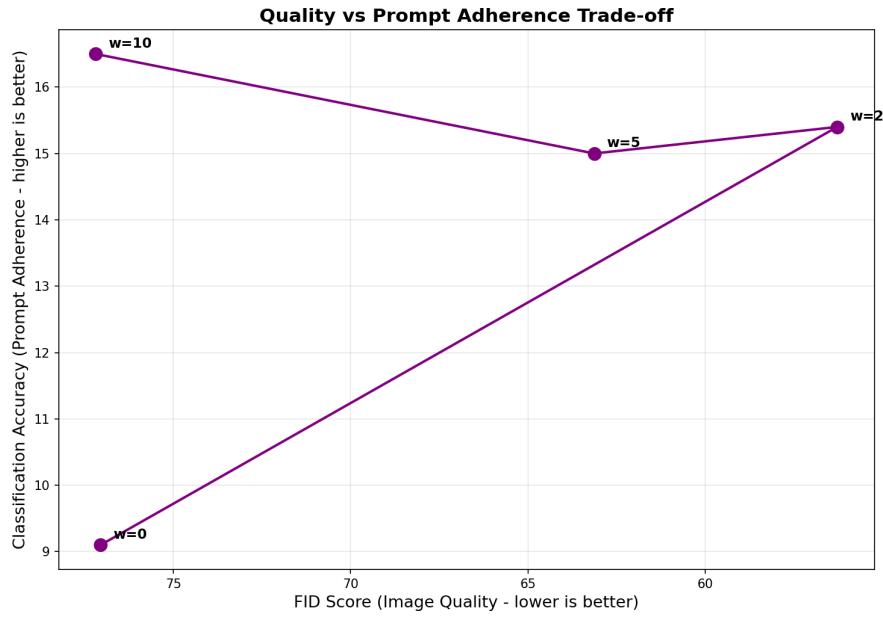


Figure 8: Quality vs. Prompt Adherence trade-off curve for CIFAR-10 generation. Each point represents a different guidance scale. The x-axis shows FID (inverted, so rightward is better quality), and the y-axis shows classification accuracy. The curve illustrates that increasing guidance improves prompt adherence at the cost of image quality.

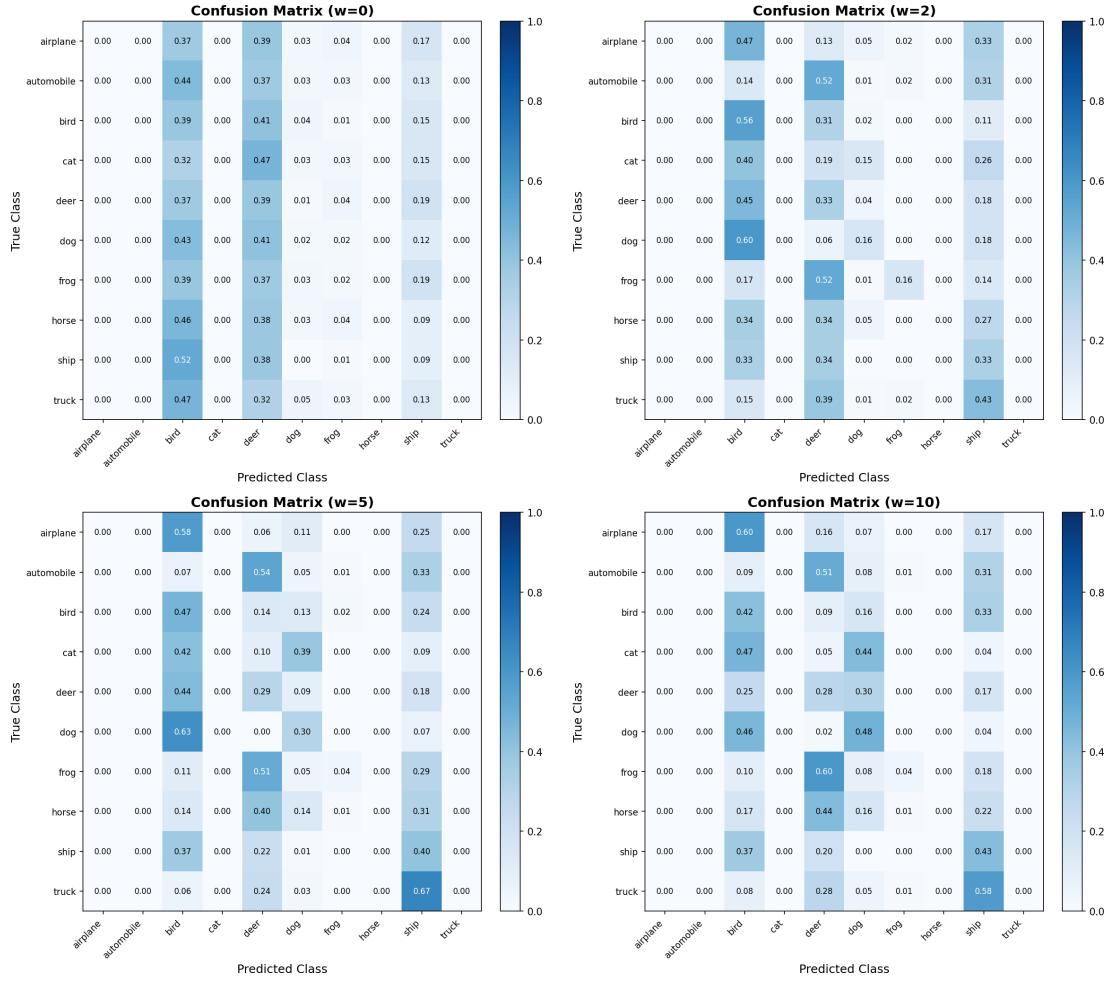


Figure 9: Normalized confusion matrices for CIFAR-10 generation across all guidance scales ($w \in \{0, 2, 5, 10\}$). Rows represent the intended class from the text prompt, columns represent the classifier's prediction. Diagonal values indicate correct generation. The matrices reveal that some classes (e.g., truck, ship) are easier to generate correctly than others (e.g., cat, dog).

Figure 10: Comparison of generated images for "A handwritten digit 0" across different guidance scales ($w = 0, 5, 10, 20, 50, 100$). Images generated with 50 denoising steps and fixed random seed 422.